ECE368: Probabilistic Reasoning
Lab 1: Classification with Multinomial and Gaussian Models

You can complete this lab in a group of two. Please provide the name and student number of both members.

Name: Bharat Bhargava	**Student Number: 1010380892**
Name: Daivik Dhar	**Student Number: 1010214260**

You should hand in: 1) A scanned .pdf version of this sheet with your answers (file size should be under 2 MB); 2) one figure for Question **1**.2.(c) and two figures for Question **2**.1.(c) in the .pdf format; and 3) two Python files classifier.py and ldaqda.py that contain your code. All these files should be uploaded to Quercus.

1 Naïve Bayes Classifier for Spam Filtering

1. (a) Write down the estimators for p_d and q_d as functions of the training data $\{\mathbf{x}_n, y_n\}, n = 1, 2, \ldots, N$ using the technique of "Laplace smoothing". (1 **pt**)

> **Idea:** Laplace smoothing avoids zero probabilities for unseen words by adding a count of 1 to every word observation. Let $S = \{n : y_n = 1\}$ be the spam emails and $H = \{n : y_n = 0\}$ be the ham emails. The estimators are:
>
> $$\hat{p}_d = \frac{1 + \sum_{n \in S} x_{n,d}}{D + \sum_{n \in S} \sum_{j=1}^{D} x_{n,j}}, \quad \hat{q}_d = \frac{1 + \sum_{n \in H} x_{n,d}}{D + \sum_{n \in H} \sum_{j=1}^{D} x_{n,j}}$$
>
> where D is the vocabulary size and $x_{n,d}$ is the count of word d in email n.

(b) Complete function learn_distributions in python file classifier.py based on the expressions you derived in part (a). (1 **pt**)

2. (a) Write down the MAP rule to decide whether $y = 1$ or $y = 0$ based on its feature vector $\mathbf{x}$ for a new email $\{\mathbf{x}, y\}$. The d-th entry of $\mathbf{x}$ is denoted by x_d. Please incorporate p_d and q_d in your expression. Please assume that $\pi = 0.5$. (1 **pt**)

> The MAP decision rule compares the posterior probabilities. We decide $y = 1$ if $P(y = 1|\mathbf{x}) > P(y = 0|\mathbf{x})$. Using Bayes' rule and the assumption that priors are equal ($\pi = 1 - \pi = 0.5$), this simplifies to comparing the likelihoods:
>
> $$P(\mathbf{x}|y = 1) \underset{y=0}{\overset{y=1}{\gtrless}} P(\mathbf{x}|y = 0)$$
>
> Substituting the Multinomial model distributions and taking the natural logarithm (to prevent numerical underflow), we get the final decision rule:
>
> $$\sum_{d=1}^{D} x_d \ln(p_d) \underset{y=0}{\overset{y=1}{\gtrless}} \sum_{d=1}^{D} x_d \ln(q_d)$$

(b) Complete function classify_new_email in classifier.py, and test the classifier on the testing set. The number of Type 1 errors is [   ], and the number of Type 2 errors is [   ]. (1 **pt**)

(c) Write down the modified decision rule in the classifier such that these two types of error can be traded off. Please introduce a new parameter to achieve such a trade-off. (0.5 **pt**)

Write your code in file classifier.py to implement your modified decision rule. Test it on the testing set and plot a figure to show the trade-off between Type 1 error and Type 2 error. In the figure, the x-axis should be the number of Type 1 errors and the y-axis should be the number of Type 2 errors. Plot at least 10 points corresponding to different pairs of these two types of error in your figure. The two end points of the plot should be: 1) the point with zero Type 1 error; and 2) the point with zero Type 2 error. Please save the figure with name **nbc.pdf**. (1 **pt**)

3. Why do we need Laplace smoothing? Briefly explain what would go wrong if we do use the maximum likelihood estimators in the training process. (0.5 **pt**)